
CS317

Information Retrieval

Week 02

Muhammad Rafi

February 03, 2020

Term Vocabulary & Posting
Lists

Boolean Model

- Information need has to be translated into a Boolean expression which most users find awkward
- The Boolean queries formulated by the users are most often too simplistic
- The Boolean model imposes a binary criterion for deciding relevance
- The question of how to extend the Boolean model to accomodate partial matching and a ranking has attracted considerable attention in the past
- Two extensions of boolean model:
 - Extended Boolean Model
 - Fuzzy Set Model

Extended Boolean Model

- Proximity Search
- Ranked Retrieval
- Example
 - WestLaw

Westlaw – Commercial Systems

- Largest commercial legal search service in terms of the number of paying subscribers
- Over half a million subscribers performing millions of searches a day over tens of terabytes of text data
- The service was started in 1975.
- In 2005, Boolean search (called “Terms and Connectors” by Westlaw) was still the default, and used by a large percentage of users . . .
- . . . although ranked retrieval has been available since 1992.

Westlaw – Commercial Systems

- Information need: Information on the legal theories involved in preventing the disclosure of trade secrets by employees formerly employed by a competing company
- Query: “trade secret” /s disclos! /s prevent /s employe!

Boolean Retrieval Model

- Last Chapter: Simple Boolean retrieval system
- Our assumptions were:
 - We know what a document is.
 - We can “machine-read” each document.
- This can be complex in reality.

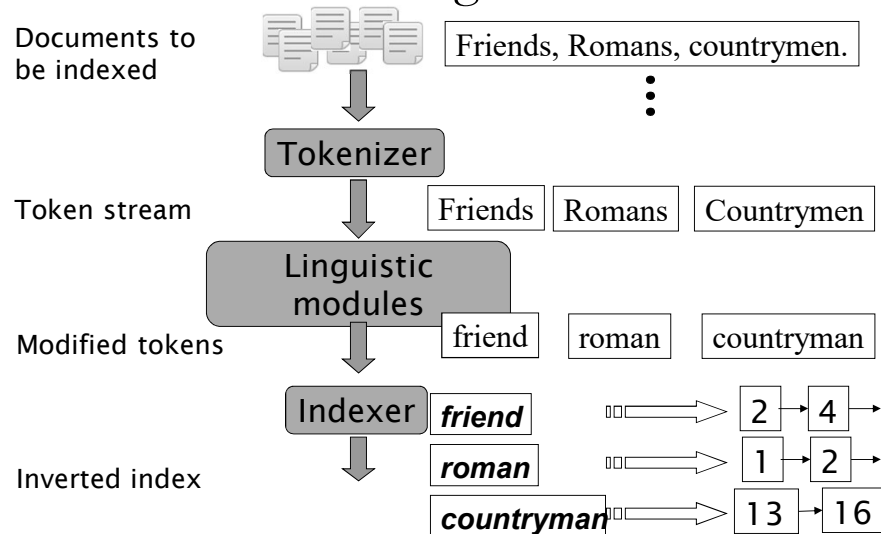
Agenda

- Document Processing

Some Definitions

- Word – A delimited string of characters as it appears in the text.
- Term – A “normalized” word (case, morphology, spelling etc); an equivalence class of words.
- Token – An instance of a word or term occurring in a document.
- Type – The same as a term in most cases: an equivalence class of tokens.

Document Processing



Challenges in Document Processing

- What format is it in?
 - pdf/word/excel/html?
- What language is it in?
- What character set is in use?
 - (CP1252, UTF-8, ...)
- Format/Language/Encoding...
- Each of these is a classification problem, which we will study later in the course.

Challenges in Document Processing

- Documents –a general term for IR
- Size of document
 - A file / An e-mail / A blog
 - A group of files.
- Tokenization
 - A process through with documents are parsed and a sequence of characters separated, as a feature for document processing.
 - A token is an instance of a sequence of characters

Challenges in Document Processing

■ Tokenization

- Tokenization process decide when to emit a token.
- Input: “Friends, Romans and Countrymen”
- Output: Tokens
 - Friends
 - Romans
 - Countrymen

■ Issues in Tokenization

- Finland’s capital

Challenges in Document Processing

■ Issues in Tokenization

- Finland’s capital
- Hewlett-Packard
- co-education
- San Francisco: one token or two?
- Numbers
 - 3/20/91
 - Mar. 12, 1991
 - 55 B.C.
 - B-52
 - (800) 234-2333

Challenges in Document Processing

■ Issues in Tokenization

□ Languages

- French
- German
- Urdu & Arabic
- Korean
- Chinese & Japanese

莎拉波娃现在居住在美国东南部的佛罗里达。今年4月9日，莎拉波娃在美国第一大城市纽约度过了18岁生日。生日派对上，莎拉波娃露出了甜美的微笑。

استقلت الجزائر في سنة 1962 بعد 132 عاما من الاحتلال الفرنسي.

← → ← → ← START

‘Algeria achieved its independence in 1962 after 132 years of French occupation.’

Sec. 2.2.2

Stop words

- With a stop list, you exclude from the dictionary entirely the commonest words. Intuition:
 - They have little semantic content: *the, a, and, to, be*
 - There are a lot of them: ~30% of postings for top 30 words
- But the trend is away from doing this:
 - Good compression techniques (lecture 5) means the space for including stopwords in a system is very small
 - Good query optimization techniques (lecture 7) mean you pay little at query time for including stop words.
 - You need them for:
 - Phrase queries: “King of Denmark”
 - Various song titles, etc.: “Let it be”, “To be or not to be”
 - “Relational” queries: “flights to London”

Sec. 2.2.3

Normalization to terms

- We need to “normalize” words in indexed text as well as query words into the same form
 - We want to match **U.S.A.** and **USA**
 - Result is terms: a term is a (normalized) word type, which is an entry in our IR system dictionary
 - We most commonly implicitly define equivalence classes of terms by, e.g.,
 - deleting periods to form a term
 - **U.S.A., USA**
 - deleting hyphens to form a term
 - **anti-discriminatory, antidiscriminatory antidiscriminatory**
-

Sec. 2.2.3

Normalization: other languages

- Accents: e.g., French **résumé** vs. **resume**.
 - Umlauts: e.g., German: **Tuebingen** vs. **Tübingen**
 - Cedilla/diacritic
 - Most important criterion:
 - How are your users like to write their queries for these words?
 - Even in languages that standardly have accents, users often may not type them
 - Often best to normalize to a de-accented term
 - **Tuebingen, Tübingen, Tubingen \ Tubingen**
-

Morphological Analysis

- Inflections: adding a suffix to a word, that doesn't change its grammatical category, such as tenses in verbs (-ing, -ed, -s), plural in nouns (s).
 - Derivations - adding a suffix to a word, that changes its grammatical category, such as nation (noun) => national (adjective) => nationalize (verb).
-

Sec. 2.2.3

Case folding

- Reduce all letters to lower case
 - exception: upper case in mid-sentence?
 - e.g., General Motors
 - Fed vs. fed
 - SAIL vs. sail
 - Often best to lower case everything, since users will use lowercase regardless of 'correct' capitalization...
-

Sec. 2.2.3

Normalization to terms

- An alternative to equivalence classing is to do asymmetric expansion
 - An example of where this may be useful
 - Enter: **window** Search: **window, windows**
 - Enter: **windows** Search: **Windows, windows, window**
 - Enter: **Windows** Search: **Windows**
 - Potentially more powerful, but less efficient
-

Thesauri and soundex

- Do we handle synonyms and homonyms?
 - E.g., by hand-constructed equivalence classes
 - **car = automobile color = colour**
 - We can rewrite to form equivalence-class terms
 - When the document contains **automobile**, index it under **car-automobile** (and vice-versa)
 - Or we can expand a query
 - When the query contains **automobile**, look under **car** as well
 - What about spelling mistakes?
 - One approach is Soundex, which forms equivalence classes of words based on phonetic heuristics
-

Lemmatization

- Lemmatization implies doing “proper” reduction to dictionary headword form
- Reduce inflectional/variant forms to base form
- E.g.,
 - *am, are, is* → *be*
 - *car, cars, car's, cars'* → *car*
- *the boy's cars are different colors* → *the boy car be different color*

Lemmatization Vs. Stemming

- Lemmatization is the algorithmic process of determining the lemma for a given word.
 - The process may involve complex tasks such as
 - understanding context and determining the part of speech of a word in a sentence (requiring, for example, knowledge of the grammar of a language)
 - it can be a hard task to implement a lemmatizer for a new language.
 - For example, in English, the verb 'to walk' may appear as 'walk', 'walked', 'walks', 'walking'. The base form, 'walk', that one might look up in a dictionary, is called the lemma for the word.

Lemmatization Vs. Stemming

- Stemmer:
 - The difference between lemmatization and stemming is that a stemmer operates on a single word **without knowledge of the context**
 - Stemmers are typically easier to implement and run faster, and the reduced accuracy may not matter for some applications

Porter Stemmer

- An incoming word is cleaned up in the initialization phase, one prefix trimming phase then takes place and then five suffix trimming phases occur.
- Note: The entire algorithm will not be covered -- we will leave out some obscure rules.

Initialization

- First the word is cleaned up. Converted to lower case only letters or digits are kept.
- F-16 is converted to f16.

Porter Stemming

- Remove prefixes:
 - "kilo", "micro", "milli", "intra", "ultra", "mega", "nano", "pico", "pseudo"
- So megabyte, kilobyte all become "byte".

Porter Step 1

- Examples:
- Remove “es” from words that end in “sses” or “ies”
 - passes --> pass, cries --> cri
- Remove “s” from words whose next to last letter is not an “s”
 - runs --> run, fuss --> fuss
- If word has a vowel and ends with “eed” remove the “ed”
 - agreed --> agre, freed --> freed
- Replace trailing “y” with an “i” if word has a vowel
 - satisfy --> satisfi, fly --> fli

Porter Step 2

- With what is left, replace any suffix on the left with suffix on the right ...
- | | | |
|-----------|------|---------------------------------|
| ■ tional | tion | conditional --> condition |
| ■ ization | ize | nationalization --> nationalize |
| ■ iveness | ive | effectiveness --> effective |
| ■ fulness | ful | usefulness --> useful |
| ■ ousness | ous | nervousness --> nervous |
| ■ ousli | ous | nervously --> nervous |
| ■ entli | ent | fervently --> fervent |
| ■ iveness | ive | inventiveness --> inventive |
| ■ biliti | ble | sensibility --> sensible |

Porter Step 3

- With what is left, replace any suffix on the left with suffix on the right ...
- icate ic fabricate --> fabric (*Think about this one*)
- ative -- combativ --> comb (*another good one*)
- alize al nationalize --> national
- iciti ic
- ical ic tropical --> tropic
- ful -- faithful --> faith
- iveness ive inventiveness --> inventive
- ness -- harness --> har

Porter Step 4

- Remove remaining standard suffixes
 - *al, ance, ence, er, ic, able, ible, ant, ement, ment, ent, sion, tion, ou, ism, ate, iti, ous, ive, ize, ise*

Porter Step 5

- Remove trailing “e” if word does not end in a vowel
 - hinge --> hing
 - free --> free

Porter Stemmer: Experimental Results

Suffix stripping of a vocabulary of 10,000 words

Number of words reduced in	step 1:	3597
	step 2:	766
	step 3:	327
	step 4:	2424
	step 5:	1373
Number of words not reduced:		3650

The resulting vocabulary of stems contained 6370 distinct entries. Thus the suffix stripping process reduced the size of the vocabulary by about one third.

Example

Sample text: Such an analysis can reveal features that are not easily visible from the variations in the individual genes and can lead to a picture of expression that is more biologically transparent and accessible to interpretation

Porter stemmer: such an analysi can reveal featur that ar not easili visibl from the variat in the individu gene and can lead to a pictur of express that is more biolog transpar and access to interpret

Lovins stemmer: such an analys can reve featur that ar not eas vis from th vari in th individu gen and can lead to a pictur of expres that is mor biolog transpar and acces to interpre

Paice stemmer: such an analys can rev feat that are not easy vis from the vary in the individ gen and can lead to a pict of express that is mor biolog transp and access to interpret

Porter Summary

- Do stemming and other normalizations help?
 - English: very mixed results. Helps recall but harms precision
 - operative (dentistry) ⇒ oper
 - operational (research) ⇒ oper
 - operating (systems) ⇒ oper
 - Definitely useful for Spanish, German, Finnish, ...
 - 30% performance gains for Finnish!
- Full morphological analysis – at most modest benefits for retrieval

Implementation Issues

- Inverted Index
 - Lists
 - Hashmap
 - Trees
- SkipList

Problem

- Are the following statements true or false?
 - In a Boolean retrieval system, stemming never lowers precision.
 - In a Boolean retrieval system, stemming never lowers recall.
 - Stemming increases the size of the vocabulary.
 - Stemming should be invoked at indexing time but not while processing a query.

Problem (Solution)

- Are the following statements true or false?
 - In a Boolean retrieval system, stemming never lowers precision. (False)
 - In a Boolean retrieval system, stemming never lowers recall. (True)
 - Stemming increases the size of the vocabulary. (False)
 - Stemming should be invoked at indexing time but not while processing a query. (False)